



# Data-Driven Decision Support for Aircraft Procurement

**Student:** Daniel Mwaka

**Student Pace:** DSF-FT12-Hybrid

**Instructor:** Samuel Karu

## Overview

This project aims to support the company's data-driven decision to procure relatively low-risk airplane models to operate the fleet in the commercial and private enterprise sectors. It analyzes a dataset compiled by the National Transportation Safety board and retrieved from Kaggle.com. Presented recommendations are deduced from descriptive statistics results and backed-up by interactive visualizations modeled using Tableau Public.

## Business Problem

Venturing into a new industry avails a company of potential growth opportunities. Diversifying a company's portfolio lowers reliance on a single revenue stream (asset diversification) and optimizes resilience against unfavorable business environments/ factors (Luo, 2022). However, venturing into a highly sensitive sector, such as operating airplanes for commercial and private enterprises, necessitates data-driven decisions, strategic implementation, and formative performance evaluation (Altundag & Wynn, 2024). Purchasing airplane models prone to crashing poses a substantial risk to human life, brand image, competitiveness, profitability, and the company's longevity once it ventures into the new industry.

## Data Understanding

The NTSB aviation accident database contains information on crashes and contingency incidents within the U.S., its territories, and across international waters from 1962 to 2023. The dataset was retrieved from Kaggle.com (NTSB, 2023). The dataset is formatted as a .csv file with 31 columns and 88889 rows. The columns are meticulously organized to capture variables of interest in aircraft accidents and incidents. Extracting, analyzing, and visualizing relevant data from the dataset is vital to shedding insight into the safest, low-risk airplane models for a new entrant to the commercial aviation industry.

However, the dataset is enormous, contains some missing values duplicates, and needs to be customized to meet new entrants' needs aiming to purchase and operate airplanes for commercial and private enterprises. For instance, it contains data for airplane crashes across multiple scenarios (purpose of the flight), in which some cases exceed the target scope of the company. For this project, the target variables are categorized into Dimensions and Measures. The Dimension variables include: Event Date, Investigation Type, Aircraft Damage, Location, Make, Model, Weather Condition, and Purpose of Flight. The Measure variables include: The number of Engines, Total Fatal Injuries, Total Minor Injuries, and Total Uninjured passengers. The insights yielded from analyzing these variables and visualizing their respective relationships are deemed appropriate for deducing recommendations to support data-driven decisions by the company to procure a fleet that comprises safe, low-risk airplanes

```
In [3]: # Importing standard packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
```

```
In [4]: # Loading the dataset and creating the master dataframe
df_master = pd.read_csv("Data/AviationData.csv", encoding='latin1', low
df_master.shape
print(f"This data set consists of {df_master.shape[0]} rows")
print(f"This data set consists of {df_master.shape[1]} columns")
```

```
This data set consists of 88889 rows
This data set consists of 31 columns
```

Copying the initially loaded DataFrame to perform ETL processes without modifying df\_master.

```
In [5]: df=df_master.copy()
df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

```
This data set consists of 88889 rows
This data set consists of 31 columns
```

In [6]:

df.head()

Out [6]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States	
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States	
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States	36
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States	
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States	

5 rows × 31 columns

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                   88889 non-null  object
2   Accident.Number                     88889 non-null  object
3   Event.Date                          88889 non-null  object
4   Location                            88837 non-null  object
5   Country                            88663 non-null  object
6   Latitude                            34382 non-null  object
7   Longitude                           34373 non-null  object
8   Airport.Code                        50132 non-null  object
9   Airport.Name                        52704 non-null  object
10  Injury.Severity                     87889 non-null  object
11  Aircraft.damage                     85695 non-null  object
12  Aircraft.Category                   32287 non-null  object
13  Registration.Number                 87507 non-null  object
14  Make                               88826 non-null  object
15  Model                              88797 non-null  object
16  Amateur.Built                      88787 non-null  object
17  Number.of.Engines                   82805 non-null  float64
18  Engine.Type                         81793 non-null  object
19  FAR.Description                     32023 non-null  object
20  Schedule                           12582 non-null  object
21  Purpose.of.flight                  82697 non-null  object
22  Air.carrier                         16648 non-null  object
23  Total.Fatal.Injuries                77488 non-null  float64
24  Total.Serious.Injuries              76379 non-null  float64
25  Total.Minor.Injuries                76956 non-null  float64
26  Total.Uninjured                     82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight               61724 non-null  object
29  Report.Status                       82505 non-null  object
30  Publication.Date                    75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

It is evident that the columns from 4th index to the 30th index are missing some data values.

```
In [8]: df.dtypes
```

```
Out[8]: Event.Id                object
Investigation.Type            object
Accident.Number              object
Event.Date                   object
Location                     object
Country                      object
Latitude                     object
Longitude                    object
Airport.Code                 object
Airport.Name                 object
Injury.Severity              object
Aircraft.damage              object
Aircraft.Category            object
Registration.Number          object
Make                         object
Model                       object
Amateur.Built                object
Number.of.Engines            float64
Engine.Type                  object
FAR.Description              object
Schedule                     object
Purpose.of.flight            object
Air.carrier                  object
Total.Fatal.Injuries         float64
Total.Serious.Injuries       float64
Total.Minor.Injuries         float64
Total.Uninjured              float64
Weather.Condition            object
Broad.phase.of.flight        object
Report.Status                object
Publication.Date             object
dtype: object
```

The data type for all the variables is either an object or a float. This information reveals the need for data type transformations (to be performed later).

```
In [9]: df.isna().sum()
```

```
Out[9]: Event.Id                                0
Investigation.Type                             0
Accident.Number                               0
Event.Date                                    0
Location                                       52
Country                                       226
Latitude                                     54507
Longitude                                    54516
Airport.Code                                 38757
Airport.Name                                36185
Injury.Severity                             1000
Aircraft.damage                             3194
Aircraft.Category                           56602
Registration.Number                          1382
Make                                          63
Model                                         92
Amateur.Built                               102
Number.of.Engines                           6084
Engine.Type                                 7096
FAR.Description                             56866
Schedule                                    76307
Purpose.of.flight                           6192
Air.carrier                                 72241
Total.Fatal.Injuries                        11401
Total.Serious.Injuries                      12510
Total.Minor.Injuries                       11933
Total.Uninjured                             5912
Weather.Condition                           4492
Broad.phase.of.flight                       27165
Report.Status                               6384
Publication.Date                            13771
dtype: int64
```

## Data Preparation

Digital maturity in leveraging novel data analytics is the key driver for aircraft safety. Manufacturers use these technologies to design safer aircraft (Boyd & Stolzer, 2016). On the other hand, companies operating in the airline sector base their strategic procurement plans on data-supported decisions (Altundag & Wynn, 2024). The progressive scope in which key stakeholders in the airline sector are embracing data analytics is reflected in the cumulative number of aircraft accidents and accidents recorded by the NTSB for each decade. Additionally, structural and material design advancements progressively increase aircraft safety. Thus, the first step to cleaning the data is to convert the `Event.Date` format from an object to the datetime format and slice the DataFrame to capture data for accidents and incidents from 2000 to 2023.

```
In [10]: # Converting the 'Event.Date' column to a datetime dtype
df['Event.Date'] = pd.to_datetime(df['Event.Date'], errors='coerce')
# Incorporating conditionals to select the period between 2000 and 2023
mask_2000_2023 = (df['Event.Date'].dt.year >= 2000) & (df['Event.Date']
# Applying the masks
df = df[mask_2000_2023]
```

As captured in the time-series plot below, the number of aircraft accidents and incidents has been dropping. The selected period is deemed appropriate because analyzing data for aircraft accidents and incidents before 2000 would comprise the reliability of deduced recommendations for real-world application in the 2020s.

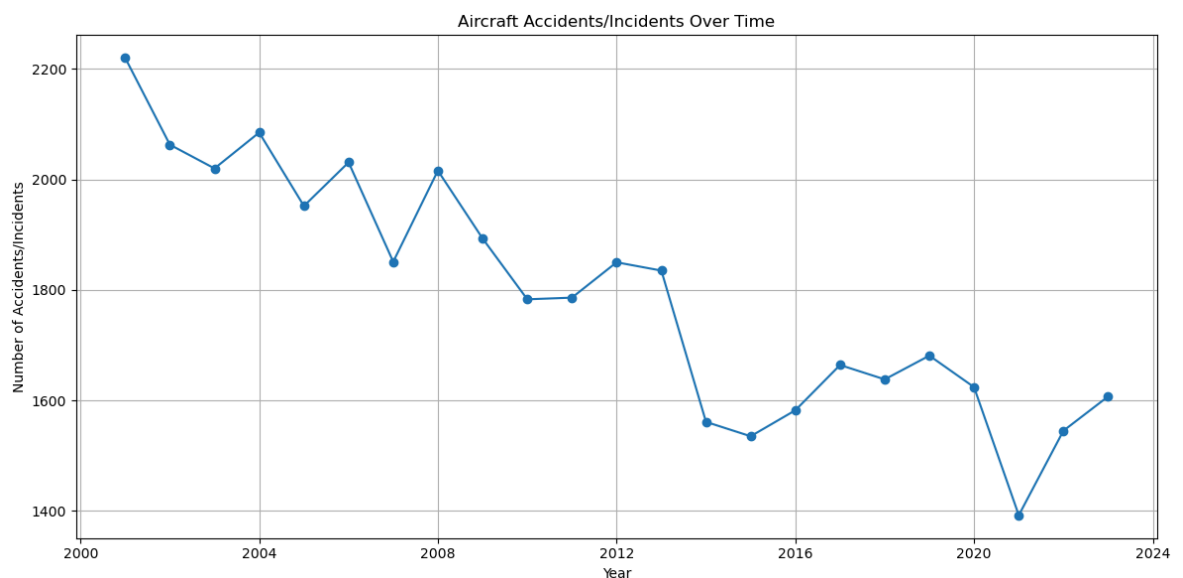
```
In [11]: # Setting the 'Event.Date' as the index
df.set_index('Event.Date', inplace=True)

# Resampling the data to count incidents per year (year-end)
yearly_counts = df.resample('Y').size()

# Creating the time series line plot
plt.figure(figsize=(12, 6))
plt.plot(yearly_counts.index, yearly_counts.values, marker='o', linestyle='solid')

plt.title('Aircraft Accidents/Incidents Over Time')
plt.xlabel('Year')
plt.ylabel('Number of Accidents/Incidents')
plt.grid(True)
plt.tight_layout()
plt.show()
```

/tmp/ipykernel\_6004/3113829321.py:5: FutureWarning: 'Y' is deprecated and will be removed in a future version, please use 'YE' instead.  
yearly\_counts = df.resample('Y').size()



The second step is dropping all the columns deemed inappropriate for this project

```
In [12]: # Dropping columns with data deemed inappropriate per the project's obj
columns_to_drop = ['Event.Id', 'Latitude', 'Longitude', 'Airport.Code',
df.drop(columns = columns_to_drop, inplace=True)
```

```
In [13]: df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

```
This data set consists of 41214 rows
This data set consists of 14 columns
```

```
In [14]: df.dtypes
```

```
Out[14]: Investigation.Type      object
Location                        object
Country                        object
Aircraft.damage                object
Make                           object
Model                          object
Number.of.Engines              float64
Engine.Type                    object
Purpose.of.flight              object
Total.Fatal.Injuries           float64
Total.Serious.Injuries         float64
Total.Minor.Injuries           float64
Total.Uninjured                float64
Weather.Condition              object
dtype: object
```

Dropping rows for entries with NaNs except for the float data type columns. The missing values for `Number.of.Engines` are also dropped because the number of engines in an aircraft despite being an interger represent an object and the variable is discrete.

```
In [15]: df = df.dropna(subset=['Location'])
df = df.dropna(subset=['Aircraft.damage'])
df = df.dropna(subset=['Make'])
df = df.dropna(subset=['Model'])
df = df.dropna(subset=['Number.of.Engines'])
df = df.dropna(subset=['Engine.Type'])
df = df.dropna(subset=['Purpose.of.flight'])
df = df.dropna(subset=['Weather.Condition'])
```



```
In [16]: df.isna().sum()
```

```
Out[16]: Investigation.Type      0
Location                        0
Country                        8
Aircraft.damage                0
Make                           0
Model                           0
Number.of.Engines              0
Engine.Type                    0
Purpose.of.flight              0
Total.Fatal.Injuries           9213
Total.Serious.Injuries         10005
Total.Minor.Injuries           9283
Total.Uninjured                4517
Weather.Condition              0
dtype: int64
```

The descriptive statistics for the float data type columns (except Number.of.Engines) are computed to determine the best approach to impute missing values.

```
In [17]: # Computing the descriptive statistics for float dtype columns
columns_to_check = ['Total.Fatal.Injuries', 'Total.Serious.Injuries', '
for col in columns_to_check:
    print(f"Descriptive Statistics for {col}:")
    print(df[col].describe())
```

```
Descriptive Statistics for Total.Fatal.Injuries:
count    20899.000000
mean      0.448251
std       1.111559
min       0.000000
25%      0.000000
50%      0.000000
75%      1.000000
max      88.000000
Name: Total.Fatal.Injuries, dtype: float64
Descriptive Statistics for Total.Serious.Injuries:
count    20107.000000
mean      0.320635
std       0.668375
min       0.000000
25%      0.000000
50%      0.000000
75%      0.000000
max       9.000000
Name: Total.Serious.Injuries, dtype: float64
Descriptive Statistics for Total.Minor.Injuries:
count    20829.000000
mean      0.305151
std       0.744433
min       0.000000
25%      0.000000
50%      0.000000
75%      0.000000
max      42.000000
Name: Total.Minor.Injuries, dtype: float64
Descriptive Statistics for Total.Uninjured:
count    25595.000000
mean      1.396992
std       5.918586
min       0.000000
25%      0.000000
50%      1.000000
75%      2.000000
max     386.000000
Name: Total.Uninjured, dtype: float64
```

All four variables exhibit right skewness, meaning most incidents have low injury counts. Since a few incidents have substantially higher counts, the median is more robust to outliers and better represents the typical value in skewed distributions. Thus, imputing missing values with the median for each column is more logical.

```
In [18]: # Imputing missing values with the median
df.loc[:, 'Total.Fatal.Injuries'] = df['Total.Fatal.Injuries'].fillna(0)
df.loc[:, 'Total.Serious.Injuries'] = df['Total.Serious.Injuries'].fillna(0)
df.loc[:, 'Total.Minor.Injuries'] = df['Total.Minor.Injuries'].fillna(0)
df.loc[:, 'Total.Uninjured'] = df['Total.Uninjured'].fillna(df['Total.Uninjured'].mean())
```

```
In [19]: df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 30112 rows  
This data set consists of 14 columns

```
In [20]: df.isna().sum()
```

```
Out[20]: Investigation.Type      0
Location                        0
Country                        8
Aircraft.damage                0
Make                           0
Model                          0
Number.of.Engines              0
Engine.Type                    0
Purpose.of.flight              0
Total.Fatal.Injuries           0
Total.Serious.Injuries         0
Total.Minor.Injuries           0
Total.Uninjured                0
Weather.Condition              0
dtype: int64
```

Although the dataset doesn't have NaNs, there could be entries assigned to an unknown variable

Using Lambda functions to drop unknown values for categorical columns

```
In [21]: df['Aircraft.damage'].value_counts()
```

```
Out[21]: Aircraft.damage
Substantial    25994
Destroyed      3732
Minor          380
Unknown         6
Name: count, dtype: int64
```

```
In [22]: #Using a lambda function to drop entries with unknown
df = df[df['Aircraft.damage'].apply(lambda which_damage: which_damage != 'Unknown')]
```

```
In [23]: df['Engine.Type'].value_counts()
```

```
Out[23]: Engine.Type
Reciprocating      26916
Turbo Prop         1367
Turbo Shaft        1338
Turbo Fan           294
Turbo Jet           145
Unknown            35
Electric            7
NONE                2
LR                  1
UNK                 1
Name: count, dtype: int64
```

```
In [24]: #Using a lambda function to drop entries with unknown
df = df[df['Engine.Type'].apply(lambda drop_unknown: (drop_unknown != '

```

```
In [25]: df['Purpose.of.flight'].value_counts()
```

```
Out[25]: Purpose.of.flight
Personal              19833
Instructional          4329
Aerial Application    1544
Business              876
Positioning           773
Other Work Use        486
Flight Test           344
Aerial Observation    325
Unknown               314
Public Aircraft       220
Ferry                 169
Executive/corporate   148
Skydiving             132
Banner Tow            94
External Load         92
Public Aircraft - Federal 86
Public Aircraft - Local  67
Public Aircraft - State  60
Air Race show         57
Air Race/show         48
Glider Tow            35
Firefighting          22
Air Drop              8
PUBS                  2
ASHO                  2
PUBL                  1
Name: count, dtype: int64
```

```
In [26]: # Using a Lambda function to select only entries whose purpose of flight
df = df[df['Purpose.of.flight'].apply(lambda niche: niche in ['Aerial A
```

```
In [27]: df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

```
This data set consists of 2568 rows
This data set consists of 14 columns
```

```
In [28]: df['Weather.Condition'].value_counts()
```

```
Out[28]: Weather.Condition
VMC      2373
IMC       191
UNK         2
Unk         2
Name: count, dtype: int64
```

```
In [29]: #Using a lambda function to drop entries with unknown
df = df[df['Weather.Condition'].apply(lambda drop_unknown: (drop_unknown
```

```
In [30]: df['Make'].value_counts()
```

```
Out[30]: Make
Cessna      265
Air Tractor  220
AIR TRACTOR INC  154
CESSNA      153
Piper       142
...
Iv Inc.      1
Curtiss-wright  1
Stinson      1
Consolidated-vultee  1
ROBINSON HELICOPTER CO  1
Name: count, Length: 294, dtype: int64
```

Converting all the values in the `Make` column to uppercase

```
In [31]: df['Make'] = df['Make'].str.upper().str.strip()
```

```
In [32]: df['Make'].value_counts()
```

```
Out[32]: Make
CESSNA      418
AIR TRACTOR  265
PIPER       231
BELL        224
AIR TRACTOR INC  156
...
WALKER      1
THRUSH AIRCRAFT INC.  1
WSK-PZL MIELIC  1
NAVION      1
ROBINSON HELICOPTER CO  1
Name: count, Length: 230, dtype: int64
```

Since there is another USState.csv file in the downloaded Zipped data from Kaggle (Presumed to be utilized in plotting a regional map in Tableau), the `Country` column is sliced to only feature rows whose value is United States

```
In [33]: # Using a lambda function to select entries for accidents and incidents
df = df[df['Country'].apply(lambda which_country: which_country == 'Uni
```

Splitting the state abbreviation section from the location's values and creating a new column `Abbreviation` to hold them. The created new column will facilitate the establishment of a relationship with the USState.csv dataset when plotting visualizations in Tableau Desktop.

```
In [34]: # Creating a new column 'Abbreviation' and extracting the Abbreviations
df['Abbreviation'] = df['Location'].apply(lambda x: x.split(',')[1] if
# Overwriting the 'Location' column with values that don't feature the A
df['Location'] = df['Location'].apply(lambda x: x.split(',')[0] if isi
# Removing the 'Abbreviation' column from the dataframe
abbreviation_col = df.pop('Abbreviation')
# Inserting the 'Abbreviation' column next to the 'Location' column
df.insert(df.columns.get_loc('Location') + 1, 'Abbreviation', abbreviat
```

```
In [ ]: # Examining whether the new column was successfully created and position
df.head()
```

```
Out[35]:
```

	Investigation.Type	Location	Abbreviation	Country	Aircraft.damage	Make
Event.Date						
2000-01-13	Accident	FILLMORE	UT	United States	Substantial	BEECH
2000-01-18	Accident	BRAWLEY	CA	United States	Substantial	BELL
2000-01-18	Accident	SOMERSET	KY	United States	Destroyed	BEECH
2000-01-20	Accident	PLAINVILLE	CT	United States	Substantial	CESSNA
2000-01-25	Accident	RAYVILLE	LA	United States	Substantial	AIR TRACTOR

Checking if there are missing values in the newly created `Abbreviations` column.

```
In [73]: # Checking the number of null values in the newly created Abbreviations
df['Abbreviation'].isna().sum()
```

```
Out[73]: np.int64(0)
```

```
In [72]: # Dropping entries that are missing values in the Abbreviation column
df = df.dropna(subset=['Abbreviation'])
```

```
In [38]: df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 2531 rows  
This data set consists of 15 columns

Checking for duplicate rows

```
In [71]: # Checking the number of duplicate entries in the DataFrame
df.duplicated().sum()
```

Out[71]: np.int64(16)

```
In [ ]: # Dropping duplicate entries
df.drop_duplicates
df.head()
```

Out[70]:

Event.Date	Investigation.Type	Location	Abbreviation	Country	Aircraft.damage	Make
2000-01-13	Accident	FILLMORE	UT	United States	Substantial	BEECH
2000-01-18	Accident	BRAWLEY	CA	United States	Substantial	BELL
2000-01-18	Accident	SOMERSET	KY	United States	Destroyed	BEECH
2000-01-20	Accident	PLAINVILLE	CT	United States	Substantial	CESSNA
2000-01-25	Accident	RAYVILLE	LA	United States	Substantial	AIR TRACTOR

```
In [69]: # Checking the data types for selected columns of interest for this pro
df.dtypes
```

Out[69]:

Investigation.Type	category
Location	object
Abbreviation	object
Country	object
Aircraft.damage	category
Make	object
Model	object
Number.of.Engines	object
Engine.Type	category
Purpose.of.flight	category
Total.Fatal.Injuries	float64
Total.Serious.Injuries	float64
Total.Minor.Injuries	float64
Total.Uninjured	float64
Weather.Condition	category
dtype:	object

Making necessary transformations for the data type of the variables to their respective appropriate d-types

```
In [42]: df['Investigation.Type'] = df['Investigation.Type'].astype('category')
df['Aircraft.damage'] = df['Aircraft.damage'].astype('category')
df['Number.of.Engines'] = df['Number.of.Engines'].astype(str)
df['Engine.Type'] = df['Engine.Type'].astype('category')
df['Purpose.of.flight'] = df['Purpose.of.flight'].astype('category')
df['Weather.Condition'] = df['Weather.Condition'].astype('category')
```

```
In [68]: # Examining whether the data type transformations were successful
df.dtypes
```

```
Out[68]: Investigation.Type      category
Location                       object
Abbreviation                   object
Country                       object
Aircraft.damage               category
Make                          object
Model                         object
Number.of.Engines             object
Engine.Type                   category
Purpose.of.flight            category
Total.Fatal.Injuries          float64
Total.Serious.Injuries        float64
Total.Minor.Injuries          float64
Total.Uninjured               float64
Weather.Condition             category
dtype: object
```

Exporting the cleaned dataset to a new .csv file

```
In [44]: df.to_csv("Data/bestest_aviation_data.csv", index=False, encoding='latin1')
```

## Data Modeling

Loading the .csv file of the cleaned data



```
In [45]: # Reading the cleaned .csv file and creating a new dataframe
df_clean = pd.read_csv("Data/bestest_aviation_data.csv",encoding='latin1')
df_clean.head()
```

```
Out[45]:
```

	Investigation.Type	Location	Abbreviation	Country	Aircraft.damage	Make	Model	Number.of.Engines
0	Accident	FILLMORE	UT	United States	Substantial	BEECH	K35	1
1	Accident	BRAWLEY	CA	United States	Substantial	BELL	OH-58C	1
2	Accident	SOMERSET	KY	United States	Destroyed	BEECH	C-90	1
3	Accident	PLAINVILLE	CT	United States	Substantial	CESSNA	T310R	1
4	Accident	RAYVILLE	LA	United States	Substantial	AIR TRACTOR	AT-401	1

```
In [67]: # Examining the columns of the df_clean DataFrame
df_clean.columns
```


```
Out[67]: Index(['Investigation.Type', 'Location', 'Abbreviation', 'Country',
                'Aircraft.damage', 'Make', 'Model', 'Number.of.Engines', 'Engine.Type',
                'Purpose.of.flight', 'Total.Fatal.Injuries', 'Total.Serious.Injuries',
                'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition'],
              dtype='object')
```

## The Least Safe Aircrafts Overall

To gain insight on the least safe aircrafts, I group the `Model` and the `Make` variable and plot a barplot against `Total.Fatal.Injuries`

```
In [66]: # Groupby 'Make' and 'Model', and sum 'Total.Uninjured'
uninjured_by_make_model = df_clean.groupby(['Make', 'Model'])['Total.Fa
# Creating a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in uninjured_b
# Creating the horizontal bar plot using seaborn
plt.figure(figsize=(10, 8))
sns.barplot(x=uninjured_by_make_model.values, y=make_model_labels, orie

plt.title('Most Risky Aircraft by Total Fatalities')
plt.xlabel('Total Fatalities')
plt.ylabel('Make and Model')
plt.tight_layout()
plt.show()
```




## The Most Safe Aircrafts Overall

To gain insight on the safest aircraft model and make, I group the `Model` and the `Make` variable and plot a barplot against `Total.Uninjured`

```
In [65]: # Grouping by 'Make' and 'Model', and sum 'Total.Uninjured'
uninjured_by_make_model = df_clean.groupby(['Make', 'Model'])['Total.Uninjured'].sum()

# Creating a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in uninjured_by_make_model.items()]

# Creating the horizontal bar plot using seaborn
plt.figure(figsize=(10, 8))
sns.barplot(x=uninjured_by_make_model.values, y=make_model_labels, orient='horizontal')
plt.title('Most Safe Aircraft by Total Uninjured Passengers')
plt.xlabel('Total Uninjured Passengers')
plt.ylabel('Make and Model')
plt.tight_layout()
plt.show()
```



To determine the safest aircraft models for each of the three civil aviation services the company can venture into; the three categorical values for the Purpose of Flight columns are plotted in a barplot against uninjured passengers Total.Uninjured .

## Recommended Aircrafts for Targeted Aviation Services' Niche

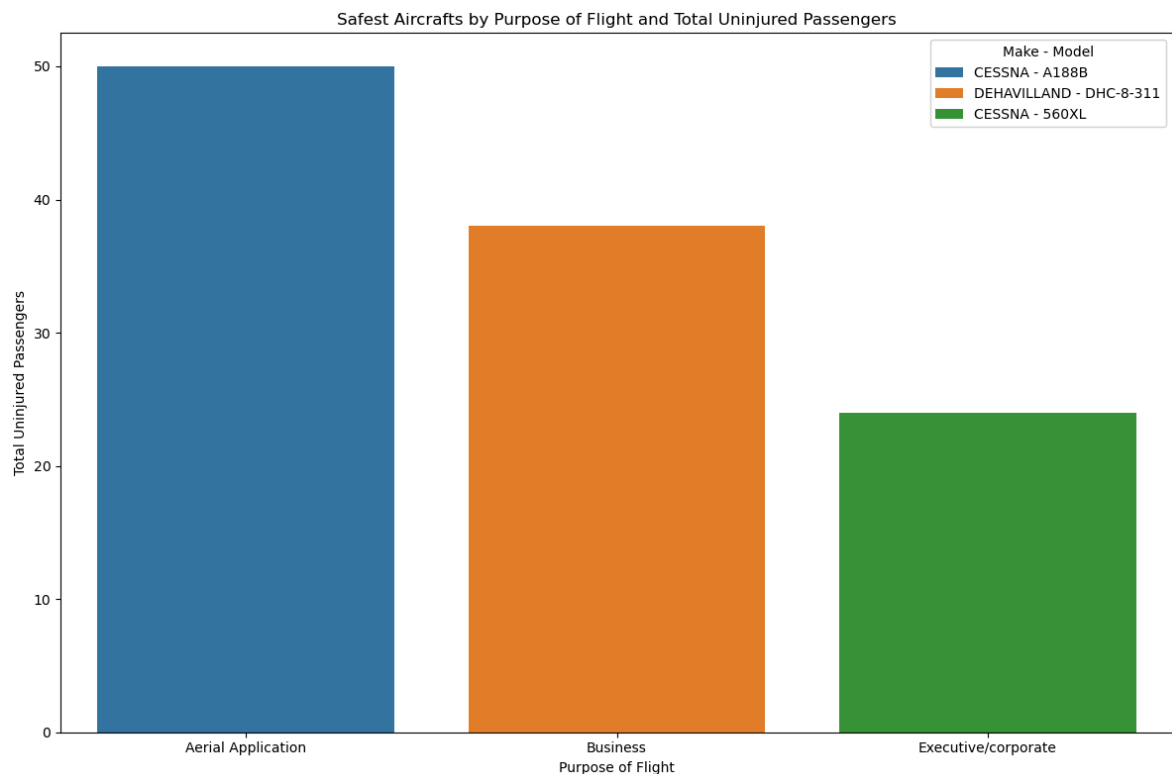
```
In [49]: # Filtering for the relevant Purpose.of.flight values
df_filtered = df_clean[df_clean['Purpose.of.flight'].isin(['Aerial Appl

# Group by 'Purpose.of.flight', 'Make', and 'Model', and sum 'Total.Uni
uninjured_by_purpose_make_model = df_filtered.groupby(['Purpose.of.flig

# Finding the safest aircraft (highest 'Total.Uninjured') for each purp
safest_aircraft = uninjured_by_purpose_make_model.loc[uninjured_by_purp

# Creating the "Make - Model" column
safest_aircraft['Make - Model'] = safest_aircraft['Make'] + ' - ' + saf

# Creating the bar plot using seaborn
plt.figure(figsize=(12, 8))
sns.barplot(x='Purpose.of.flight', y='Total.Uninjured', hue='Make - Mod
plt.title('Safest Aircrafts by Purpose of Flight and Total Uninjured Pa
plt.xlabel('Purpose of Flight')
plt.ylabel('Total Uninjured Passengers')
plt.legend(title='Make - Model')
plt.tight_layout()
plt.show()
```



## Evaluation

The baseline model shed the following insights:

- The BEECH-200 (8-seater), the CESSNA-340A (6-seater), and the BEECH-A36 (6-seater), are the top-three most risky aircraft overall.
- The CESSNA-A188B (1-seater), the BELL-206B (5-seater), and the GRUMMAN-G-164A (1-seater) are the top-three safest aircraft models overall.

### Recommendations:

- The CESSNA-560XL (10-seater) is the safest aircraft for Executive/corporate flights.
- The DEHAVILLAND-DHC-8-311 (50-seater) is the safest airplane for business flights.
- The CESSNA-A188B (1-seater) is the safest aircraft for Aerial Applications.

The number of engines for these aircraft models.

- BEECH-200: 2 engines
- CESSNA-340A: 2 engines
- BEECH-A36: 1 engine
- BELL-206B: 1 engine
- GRUMMAN-G-164A: 1 engine
- CESSNA-560XL: 2 engines
- CESSNA-A188B: 1 engine
- DEHAVILLAND-DHC-8-311: 2 engines

Multi-engine aircraft are typically safer in comparison to single-engine airplanes (Pilot Institute, 2023). More than one engine avails redundancy to propulsion units. Pilots can recalibrate controls and continue flying if one engine stalls or malfunctions. In contrast, such an incident in a single-engine aircraft poses a substantial risk of crashing if the gliding distance falls short of a runway or a relatively flat surface for an emergency landing. Thus, I modified the baseline model to drop row entries whose `Number.of.Engines` is less than 2.

```
In [57]: # Creating a copy of the df_clean dataframe to avoid modifying the base
df_modified = df_clean.copy()
df_modified.shape
print(f"This data set consists of {df_modified.shape[0]} rows")
print(f"This data set consists of {df_modified.shape[1]} columns")
```

This data set consists of 2531 rows  
This data set consists of 15 columns

```
In [58]: # Applying a lambda function to drop entries for single-engine aircraft
df_modified = df_modified[df_modified['Number.of.Engines'].apply(lambda
```

```
In [59]: # Confirming if the modifications are imputed to the df_modified datafr
df_modified['Number.of.Engines'].value_counts()
```

```
Out[59]: Number.of.Engines
2.0      369
3.0        7
4.0        3
Name: count, dtype: int64
```

```
In [60]: df_modified.shape  
print(f"This data set consists of {df_modified.shape[0]} rows")  
print(f"This data set consists of {df_modified.shape[1]} columns")  
  
This data set consists of 379 rows  
This data set consists of 15 columns
```


Replotting bar plots to determine the most risky, the safest, and the appropriate airplanes (that have more than one engine) for the company's operations

## The Least Safe Multi-Engine Aircrafts

```
In [61]: # Grouping by 'Make' and 'Model', and sum 'Total.Uninjured'
uninjured_by_make_model = df_modified.groupby(['Make', 'Model'])['Total

# Creating a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in uninjured_b

# Creating the horizontal bar plot using seaborn
plt.figure(figsize=(10, 8))
sns.barplot(x=uninjured_by_make_model.values, y=make_model_labels, orie
plt.title('Most Risky Multi-Engine Aircraft by Total Fatalities')
plt.xlabel('Total Fatalities')
plt.ylabel('Make and Model')
plt.tight_layout()
plt.show()
```




## The Most Safe Multi-Engine Aircrafts

```
In [62]: # Grouping by 'Make' and 'Model', and sum 'Total.Uninjured'
uninjured_by_make_model = df_modified.groupby(['Make', 'Model'])['Total

# Creating a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in uninjured_b

# Creating the horizontal bar plot using seaborn
plt.figure(figsize=(10, 8))
sns.barplot(x=uninjured_by_make_model.values, y=make_model_labels, orie
plt.title('Most Safe Multi-Engine Aircraft by Total Uninjured Passenger
plt.xlabel('Total Uninjured Passengers')
plt.ylabel('Make and Model')
plt.tight_layout()
plt.show()
```





## Recommended Multi-Engine Aircrafts for Targeted Aviation Services' Niche


```
In [63]: # Filtering for the relevant Purpose.of.flight values
df_filtered = df_modified[df_modified['Purpose.of.flight'].isin(['Aerial', 'Business', 'Executive/corporate'])]

# Grouping by 'Purpose.of.flight', 'Make', and 'Model', and sum 'Total.Uninjured'
uninjured_by_purpose_make_model = df_filtered.groupby(['Purpose.of.flight', 'Make', 'Model']).sum()

# Finding the safest aircraft (highest 'Total.Uninjured') for each purpose
safest_aircraft = uninjured_by_purpose_make_model.loc[uninjured_by_purpose_make_model.groupby('Purpose.of.flight')['Total.Uninjured'].idxmax()]

# Creating the "Make - Model" column
safest_aircraft['Make - Model'] = safest_aircraft['Make'] + ' - ' + safest_aircraft['Model']

# Creating the bar plot using seaborn
plt.figure(figsize=(12, 8))
sns.barplot(x='Purpose.of.flight', y='Total.Uninjured', hue='Make - Model')
plt.title('Safest Aircraft by Purpose of Flight and Total Uninjured Passengers')
plt.xlabel('Purpose of Flight')
plt.ylabel('Total Uninjured Passengers')
plt.legend(title='Make - Model')
plt.tight_layout()
plt.show()
```



The modified model shed the following insights:

- The BEECH-200 (13-seater), the CESSNA-340A (6-seater), and the BEECH-A400 (8-seater) are the top-three most risky multi-engine aircraft to operate.

- The DEHAVILLAND DHC-8-311 (50-seater), the GULFSTREAM AEROSPACE-G-IV (10-seater), and the BEECH-200 (13-seater) are the top-three safest multi-engine aircraft to operate.

### Recommendations:

- The CESSNA 560XL(10-seater) is the safest multi-engine aircraft for Executive/corporate flights.
- The DEHAVILLAND -DHC-8-311(50-seater) is the safest multi-engine aircraft for business flights.
- The MCDONNELL DOUGLAS-DC 10-10(250-seater) is the safest multi-engine aircraft for Aerial Applications.

The number of engines for these aircraft models.

- BEECH-200: 2 engines
- CESSNA-340A: 2 engines
- BEECH-A400: 2 engines
- GULFSTREAM AEROSPACE-G-IV: 2 engines
- CESSNA 560XL: 2 engines
- DEHAVILLAND DHC-8-311: 2 engines
- MCDONNELL DOUGLAS -DC 10-10: 3 engines

The findings from the results yielded by the baseline model are compared to those yielded by the modified model. After entering the industry, the recommended safest aircraft model for the aviation services the company will offer is based on respective applicability and operational logistics.

## Conclusions

The analysis yields three recommendations on the aircraft models the company should procure and operate after entering the commercial aviation industry.

- **The CESSNA-560XL aircraft is recommended for Executive/ Corporate flights:** The baseline and modified model conform the aircraft is safest for executive and corporate flights.
- **The DEHAVILLAND DHC-8-311 aircraft is recommended for Business Flights:** The baseline and modified model conform the aircraft is safest for business flights.
- **The CESSNA-A188B aircraft is recommended for Aerial Applications:** The modified model proposes the MCDONNELL DOUGLAS-DC 10-10 (a three-engine, 250-seater) aircraft for aerial applications. The proposed alternative by the modified model is rejected because aerial applications typically include agricultural activities such as spraying crop fields. Hence, the single-engine CESSNA-A188B is recommended for aerial applications.

## References

Altundag, A., & Wynn, M. (2024). Advanced Analytics and Data Management in the Procurement Function: An Aviation Industry Case Study. *Electronics*, 13(8), 1554. <https://doi.org/10.3390/electronics13081554>  
(<https://doi.org/10.3390/electronics13081554>)

Boyd, D. D., & Stolzer, A. (2016). Accident-precipitating factors for crashes in turbine-powered general aviation aircraft. *Accident Analysis & Prevention*, 86, 209-216. <https://doi.org/10.1016/j.aap.2015.10.024> (<https://doi.org/10.1016/j.aap.2015.10.024>)

NTSB. (2023). Aviation Accident Database & Synopses, up to 2023. <https://doi.org/10.34740/KAGGLE/DSV/4875576>  
(<https://doi.org/10.34740/KAGGLE/DSV/4875576>)

Luo, J. (2022). Data-driven innovation: What is it?. *IEEE Transactions on Engineering Management*, 70(2), 784-790. <https://doi.org/10.1109/TEM.2022.3145231>  
(<https://doi.org/10.1109/TEM.2022.3145231>)

Pilot Institute. (2023). Single Engine Vs. Multi Engine: Which is Better? <https://pilotinstitute.com/single-vs-multi-engine/> (<https://pilotinstitute.com/single-vs-multi-engine/>)