



# Data-Driven Decision Support for Aircraft Procurement

**Student:** Daniel Mwaka

**Student Pace:** DSF-FT12-Hybrid

**Instructor:** Samuel Karu

## Overview

This project aims to support the company's data-driven decision to procure relatively low-risk airplane models to operate the fleet in the commercial and private enterprise sectors. It analyzes a dataset compiled by the National Transportation Safety Board and retrieved from Kaggle.com. Presented recommendations are deduced from descriptive statistics results and backed-up by interactive visualizations modeled using Tableau Public.

## Business Problem

Venturing into a new industry avails a company of potential growth opportunities. Diversifying a company's portfolio lowers reliance on a single revenue stream (asset diversification) and optimizes resilience against unfavorable business environments/ factors (Luo, 2022). However, venturing into a highly sensitive sector, such as operating airplanes for commercial and private enterprises, necessitates data-driven decisions, strategic implementation, and formative performance evaluation (Altundag & Wynn, 2024). Purchasing airplane models prone to crashing poses a substantial risk to human life, brand image, competitiveness, profitability, and the company's longevity once it ventures into the new industry.

## Data Understanding

The NTSB aviation accident database contains information on crashes and contingency incidents within the U.S., its territories, and across international waters from 1962 to 2023. The dataset was retrieved from Kaggle.com (NTSB, 2023). The dataset is formatted as a .csv file with 31 columns and 88889 rows. The columns are meticulously organized to capture variables of interest in aircraft accidents and incidents. Extracting, analyzing, and visualizing relevant data from the dataset is vital to shedding insight into the safest, low-risk airplane models for a new entrant to the commercial aviation industry.

However, the dataset is enormous, contains some missing values duplicates, and needs to be customized to meet new entrants' needs aiming to purchase and operate airplanes for commercial and private enterprises. For instance, it contains data for airplane crashes across multiple scenarios (purpose of the flight), in which some cases exceed the target scope of the company. For this project, the target variables are categorized into Dimensions and Measures. The Dimension variables include: Event Date, Investigation Type, Aircraft Damage, Location, Make, Model, Weather Condition, and Purpose of Flight. The Measure variables include: The number of Engines, Total Fatal Injuries, Total Minor Injuries, and Total Uninjured passengers. The insights yielded from analyzing these variables and visualizing their respective relationships are deemed appropriate for deducing recommendations to support data-driven decisions by the company to procure a fleet that comprises safe, low-risk airplanes

```
In [1]: # Import standard packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: # Load the dataset and creating the master dataframe
df_master = pd.read_csv("data/aviation-data.csv", encoding='latin1', low_memory=False)
df_master.head()
```

Out[2]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States	NaN	N
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States	NaN	N
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States	36.922223	-81.8780
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States	NaN	N
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States	NaN	N

5 rows × 31 columns

Copying the initially loaded DataFrame to perform ETL processes without modifying df\_master.

```
In [3]: # Create a copy of df_master to apply data cleaning modifications
df=df_master.copy()

# Check the number of rows and columns of the DataFrame
df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 88889 rows  
This data set consists of 31 columns

```
In [4]: # Check columns names
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                   88889 non-null  object
2   Accident.Number                     88889 non-null  object
3   Event.Date                          88889 non-null  object
4   Location                            88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                           34382 non-null  object
7   Longitude                          34373 non-null  object
8   Airport.Code                        50249 non-null  object
9   Airport.Name                       52790 non-null  object
10  Injury.Severity                     87889 non-null  object
11  Aircraft.damage                     85695 non-null  object
12  Aircraft.Category                   32287 non-null  object
13  Registration.Number                 87572 non-null  object
14  Make                               88826 non-null  object
15  Model                              88797 non-null  object
16  Amateur.Built                      88787 non-null  object
17  Number.of.Engines                   82805 non-null  float64
18  Engine.Type                         81812 non-null  object
19  FAR.Description                     32023 non-null  object
20  Schedule                           12582 non-null  object
21  Purpose.of.flight                  82697 non-null  object
22  Air.carrier                        16648 non-null  object
23  Total.Fatal.Injuries                77488 non-null  float64
24  Total.Serious.Injuries              76379 non-null  float64
25  Total.Minor.Injuries                76956 non-null  float64
26  Total.Uninjured                     82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight               61724 non-null  object
29  Report.Status                       82508 non-null  object
30  Publication.Date                    75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

It is evident that the columns from 4th index to the 30th index are missing some data values.

```
In [5]: # Show column data types  
df.dtypes
```

```
Out[5]: Event.Id                object  
Investigation.Type            object  
Accident.Number              object  
Event.Date                   object  
Location                     object  
Country                      object  
Latitude                     object  
Longitude                    object  
Airport.Code                 object  
Airport.Name                 object  
Injury.Severity              object  
Aircraft.damage              object  
Aircraft.Category            object  
Registration.Number          object  
Make                         object  
Model                        object  
Amateur.Built                object  
Number.of.Engines            float64  
Engine.Type                  object  
FAR.Description              object  
Schedule                     object  
Purpose.of.flight            object  
Air.carrier                  object  
Total.Fatal.Injuries          float64  
Total.Serious.Injuries        float64  
Total.Minor.Injuries          float64  
Total.Uninjured              float64  
Weather.Condition            object  
Broad.phase.of.flight        object  
Report.Status                 object  
Publication.Date             object  
dtype: object
```

The data type for all the variables is either an object or a float. This information reveals the need for data type transformations (to be performed later).

```
In [6]: # Print the count of missing values for each column
df.isna().sum()
```

```
Out[6]: Event.Id                                0
Investigation.Type                             0
Accident.Number                               0
Event.Date                                    0
Location                                       52
Country                                       226
Latitude                                     54507
Longitude                                    54516
Airport.Code                                 38640
Airport.Name                                 36099
Injury.Severity                             1000
Aircraft.damage                             3194
Aircraft.Category                           56602
Registration.Number                          1317
Make                                          63
Model                                         92
Amateur.Built                               102
Number.of.Engines                           6084
Engine.Type                                 7077
FAR.Description                             56866
Schedule                                    76307
Purpose.of.flight                           6192
Air.carrier                                 72241
Total.Fatal.Injuries                        11401
Total.Serious.Injuries                     12510
Total.Minor.Injuries                       11933
Total.Uninjured                             5912
Weather.Condition                           4492
Broad.phase.of.flight                       27165
Report.Status                               6381
Publication.Date                            13771
dtype: int64
```

## Data Preparation

Digital maturity in leveraging novel data analytics is the key driver for aircraft safety. Manufacturers use these technologies to design safer aircraft (Boyd & Stolzer, 2016). On the other hand, companies operating in the airline sector base their strategic procurement plans on data-supported decisions (Altundag & Wynn, 2024). The progressive scope in which key stakeholders in the airline sector are embracing data analytics is reflected in the cumulative number of aircraft accidents and accidents recorded by the NTSB for each decade. Additionally, structural and material design advancements progressively increase aircraft safety.

The first step in cleaning the data is to check if the dataset contains duplicate entries.

```
In [7]: # Check the number of duplicate entries in the DataFrame
df.duplicated().sum()
```

```
Out[7]: 0
```

```
In [8]: # Drop duplicate entries
df.drop_duplicates(inplace=True)
```

```
In [9]: # Confirm if duplicates are removed
df.duplicated().sum()
```

Out[9]: 0

The second step is to convert the `Event.Date` format from an object to the datetime format and slice the DataFrame to capture data for accidents and incidents from 2000 to 2023.

```
In [10]: # Convert the 'Event.Date' column to a datetime dtype
df['Event.Date'] = pd.to_datetime(df['Event.Date'])

# Incorporate conditionals to select the period between 2000 and 2023
mask_2000_2023 = (df['Event.Date'].dt.year >= 2000) & (df['Event.Date'].dt.year <= 2023)

# Apply the mask to impute the sliced df
df = df[mask_2000_2023]
```

```
In [11]: # Set the 'Event.Date' as the index
df.set_index('Event.Date', inplace=True)

# Resample the data to count incidents per year (year-end)
yearly_counts = df.resample('Y').size()

# Create a time series line plot
plt.figure(figsize=(10, 6))
plt.plot(yearly_counts.index, yearly_counts.values, marker='o', linestyle='--')

# Customize plot title, axes, and display grid for easy visibility
plt.title('Aircraft Accidents/Incidents From 2000 to 2023')
plt.xlabel('Year')
plt.ylabel('Number of Accidents/Incidents')
plt.grid(True)

plt.tight_layout()

# Save the plot to the images folder
plt.savefig("./images/time-series-plot.png", dpi=300, facecolor='white')
plt.show()
```



As captured in the time-series plot, the number of aircraft accidents and incidents has been dropping. The selected period is deemed appropriate because analyzing data for aircraft accidents and incidents before 2000 would comprise the reliability of deduced recommendations for real-world application in the 2020s. The next step is to drop all the columns deemed inappropriate for this project

```
In [12]: # Drop columns with data deemed inappropriate per the project's objectives
columns_to_drop = ['Event.Id', 'Latitude', 'Longitude', 'Airport.Code', 'Airport.Name']
df.drop(columns = columns_to_drop, inplace=True)
```

```
In [13]: # Print number of rows and columns of the DataFrame
df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 41214 rows  
This data set consists of 14 columns

```
In [14]: # Display the data types of each column
df.dtypes
```

```
Out[14]: Investigation.Type      object
Location                       object
Country                       object
Aircraft.damage               object
Make                          object
Model                         object
Number.of.Engines             float64
Engine.Type                   object
Purpose.of.flight             object
Total.Fatal.Injuries          float64
Total.Serious.Injuries        float64
Total.Minor.Injuries          float64
Total.Uninjured               float64
Weather.Condition              object
dtype: object
```

Dropping rows for entries with NaNs except for the float data type columns. The missing values for `Number.of.Engines` are also dropped because the number of engines in an aircraft despite being an interger represent an object and the variable is discrete.

```
In [15]: # Remove rows with missing values in selected categorical and discrete columns
df = df.dropna(subset=['Location'])
df = df.dropna(subset=['Aircraft.damage'])
df = df.dropna(subset=['Make'])
df = df.dropna(subset=['Model'])
df = df.dropna(subset=['Number.of.Engines'])
df = df.dropna(subset=['Engine.Type'])
df = df.dropna(subset=['Purpose.of.flight'])
df = df.dropna(subset=['Weather.Condition'])
```

```
In [17]: # Print the count of missing values for each column in the sliced DataFrame
df.isna().sum()
```

```
Out[17]: Investigation.Type      0
Location                       0
Country                       8
Aircraft.damage               0
Make                          0
Model                         0
Number.of.Engines             0
Engine.Type                   0
Purpose.of.flight             0
Total.Fatal.Injuries          9213
Total.Serious.Injuries        10005
Total.Minor.Injuries          9283
Total.Uninjured               4517
Weather.Condition              0
dtype: int64
```



The descriptive statistics for the float data type columns (except Number.ofEngines) are computed to determine the best approach to impute missing values.

```
In [66]: # Compute the descriptive statistics for float dtype columns
columns_to_check = ['Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.

# Iterating for each float dtype column over columns_to_check list
for col in columns_to_check:
    print(f"Descriptive Statistics for {col}:")
    print(df[col].describe())
```

Descriptive Statistics for Total.Fatal.Injuries:

```
count    20912.000000
mean         0.447972
std         1.111269
min         0.000000
25%         0.000000
50%         0.000000
75%         1.000000
max         88.000000
```

Name: Total.Fatal.Injuries, dtype: float64

Descriptive Statistics for Total.Serious.Injuries:

```
count    20120.000000
mean         0.320974
std         0.668653
min         0.000000
25%         0.000000
50%         0.000000
75%         0.000000
max         9.000000
```

Name: Total.Serious.Injuries, dtype: float64

Descriptive Statistics for Total.Minor.Injuries:

```
count    20842.000000
mean         0.305057
std         0.744264
min         0.000000
25%         0.000000
50%         0.000000
75%         0.000000
max         42.000000
```

Name: Total.Minor.Injuries, dtype: float64

Descriptive Statistics for Total.Uninjured:

```
count    25608.000000
mean         1.398899
std         5.919773
min         0.000000
25%         0.000000
50%         1.000000
75%         2.000000
max        386.000000
```

Name: Total.Uninjured, dtype: float64

All four variables exhibit right skewness, meaning most incidents have low injury counts. Since a few incidents have substantially higher counts, the median is more robust to outliers and better represents the typical value in skewed distributions. Thus, imputing missing values with the median for each column is more logical.

```
In [18]: # Impute missing values with medians
df.loc[:, 'Total.Fatal.Injuries'] = df['Total.Fatal.Injuries'].fillna(df['Total.Fatal.Injuries'].median())
df.loc[:, 'Total.Serious.Injuries'] = df['Total.Serious.Injuries'].fillna(df['Total.Serious.Injuries'].median())
df.loc[:, 'Total.Minor.Injuries'] = df['Total.Minor.Injuries'].fillna(df['Total.Minor.Injuries'].median())
df.loc[:, 'Total.Uninjured'] = df['Total.Uninjured'].fillna(df['Total.Uninjured'].median())
```

```
In [19]: # Check the number of rows and columns of the DataFrame
df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")

This data set consists of 30125 rows
This data set consists of 14 columns
```

```
In [20]: # Confirm no NaNs in sliced DataFrame
df.isna().sum()
```

```
Out[20]: Investigation.Type      0
Location                      0
Country                      8
Aircraft.damage              0
Make                        0
Model                      0
Number.of.Engines            0
Engine.Type                 0
Purpose.of.flight           0
Total.Fatal.Injuries        0
Total.Serious.Injuries      0
Total.Minor.Injuries        0
Total.Uninjured             0
Weather.Condition           0
dtype: int64
```

Although the dataset doesn't have NaNs, there could be entries assigned to an unknown variable

Using Lambda functions to drop unknown values for categorical columns

```
In [21]: # Use the .value_counts() method to check counts of unique values
df['Aircraft.damage'].value_counts()
```

```
Out[21]: Substantial      26006
Destroyed      3733
Minor          380
Unknown         6
Name: Aircraft.damage, dtype: int64
```

```
In [22]: # Apply a lambda function to drop entries with unknown
df = df[df['Aircraft.damage'].apply(lambda which_damage: which_damage != 'Unknown')]
```

```
In [23]: # Call the .value_counts() method to check counts of unique values
df['Engine.Type'].value_counts()
```

```
Out[23]: Reciprocating      26916
Turbo Prop      1367
Turbo Shaft      1338
Turbo Fan      294
Turbo Jet      145
Unknown      35
None      13
Electric      7
NONE      2
LR      1
UNK      1
Name: Engine.Type, dtype: int64
```

```
In [24]: # Apply a lambda function to drop entries with unknown
df = df[df['Engine.Type'].apply(lambda drop_unknown: (drop_unknown != 'Unknown') &
```

```
In [25]: # Use the .value_counts() method to check counts of unique values
df['Purpose.of.flight'].value_counts()
```

```
Out[25]: Personal      19838
Instructional      4332
Aerial Application      1544
Business      879
Positioning      773
Other Work Use      487
Flight Test      344
Aerial Observation      326
Unknown      314
Public Aircraft      220
Ferry      169
Executive/corporate      148
Skydiving      132
Banner Tow      94
External Load      92
Public Aircraft - Federal      86
Public Aircraft - Local      67
Public Aircraft - State      60
Air Race show      57
Air Race/show      48
Glider Tow      35
Firefighting      22
Air Drop      8
PUBS      2
ASHO      2
PUBL      1
Name: Purpose.of.flight, dtype: int64
```

```
In [26]: # Apply a Lambda function to select only entries whose purpose of flight are relevant
df = df[df['Purpose.of.flight'].apply(lambda niche: niche in ['Aerial Application',
```

```
In [27]: # Check the number of rows and columns of the DataFrame
df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 2571 rows  
This data set consists of 14 columns

```
In [28]: # Call the .value_counts() method to check counts of unique values
df['Weather.Condition'].value_counts()
```

```
Out[28]: VMC      2376
IMC       191
Unk        2
UNK        2
Name: Weather.Condition, dtype: int64
```

```
In [29]: # Apply a lambda function to drop entries with unknown
df = df[df['Weather.Condition'].apply(lambda drop_unknown: (drop_unknown != 'Unk'))]
```

```
In [30]: # Use the .value_counts() method to check counts of unique values
df['Make'].value_counts()
```

```
Out[30]: Cessna      265
Air Tractor    220
AIR TRACTOR INC 154
CESSNA        153
Piper         142
...
Textron Aviation Inc 1
Bell/garlick      1
GARLICK          1
Air Tractor, Inc. 1
Gulfstream       1
Name: Make, Length: 297, dtype: int64
```

Converting all the values in the `Make` column to uppercase

```
In [31]: # Use string methods to perform element wise manipulations
df['Make'] = df['Make'].str.upper().str.strip()
```

```
In [32]: # Call the value_counts() method to check counts of unique values
df['Make'].value_counts()
```

```
Out[32]: CESSNA      418
AIR TRACTOR    265
PIPER         231
BELL          224
AIR TRACTOR INC 156
...
SCOTTS-BELL 47 INC 1
NAVION           1
STEMME GMBH & CO  1
COMMANDER AIRCRAFT 1
AERO VODOCHODY   1
Name: Make, Length: 233, dtype: int64
```

Since there is another USState.csv file in the downloaded Zipped data from Kaggle (Presumed to be utilized in plotting a regional map in Tableau), the `Country` column is sliced to only feature rows whose value is United

```
In [33]: # Apply a lambda function to select entries for accidents and incidents that happen
df = df[df['Country'].apply(lambda which_country: which_country == 'United States')]
```

Splitting the state abbreviation section from the location's values and creating a new column Abbreviation to hold them. The created new column will facilitate the establishment of a relationship with the USState.csv dataset when plotting visualizations in Tableau Desktop.

```
In [34]: # Create a new column 'Abbreviation' and extracting the Abbreviations for the state
df['Abbreviation'] = df['Location'].apply(lambda x: x.split(', ')[-1] if isinstance(x, str) else '')

# Overwrite the 'Location' column with values that dont feature the Abbreviation part
df['Location'] = df['Location'].apply(lambda x: x.split(', ')[0] if isinstance(x, str) else '')

# Remove the 'Abbreviation' column from the dataframe
abbreviation_col = df.pop('Abbreviation')

# Insert the 'Abbreviation' column next to the 'Location' column
df.insert(df.columns.get_loc('Location') + 1, 'Abbreviation', abbreviation_col)
```

```
In [35]: # Examine whether the new column was successfully created and positioned adjacent to the Location column
df.head()
```

Out[35]:

	Investigation.Type	Location	Abbreviation	Country	Aircraft.damage	Make	Model	Number.of
Event.Date								
2000-01-13	Accident	FILLMORE	UT	United States	Substantial	BEECH	K35	
2000-01-18	Accident	BRAWLEY	CA	United States	Substantial	BELL	OH-58C	
2000-01-18	Accident	SOMERSET	KY	United States	Destroyed	BEECH	C-90	
2000-01-20	Accident	PLAINVILLE	CT	United States	Substantial	CESSNA	T310R	
2000-01-25	Accident	RAYVILLE	LA	United States	Substantial	AIR TRACTOR	AT-401	

Checking if there are missing values in the newly created Abbreviations column.

```
In [36]: # Check the number of null values in the newly created Abbreviations column
df['Abbreviation'].isna().sum()
```

Out[36]: 4

```
In [37]: # Drop entries that are missing values in the Abbreviation column
df = df.dropna(subset=['Abbreviation'])
```

```
In [38]: # Check the number of rows and columns of the DataFrame
df.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 2534 rows  
This data set consists of 15 columns

Checking for duplicate rows

```
In [39]: # print data types of each column
df.dtypes
```

```
Out[39]: Investigation.Type      object
Location                       object
Abbreviation                   object
Country                       object
Aircraft.damage                object
Make                          object
Model                         object
Number.of.Engines              float64
Engine.Type                   object
Purpose.of.flight              object
Total.Fatal.Injuries           float64
Total.Serious.Injuries         float64
Total.Minor.Injuries           float64
Total.Uninjured                float64
Weather.Condition              object
dtype: object
```

Making necessary transformations for the data type of the variables to their respective appropriate d-types

```
In [40]: # Convert columns to appropriate data types
df['Investigation.Type'] = df['Investigation.Type'].astype('category')
df['Aircraft.damage'] = df['Aircraft.damage'].astype('category')
df['Number.of.Engines'] = df['Number.of.Engines'].astype(str)
df['Engine.Type'] = df['Engine.Type'].astype('category')
df['Purpose.of.flight'] = df['Purpose.of.flight'].astype('category')
df['Weather.Condition'] = df['Weather.Condition'].astype('category')
```

```
In [41]: # Confirm whether the data transformations are successful
df.dtypes
```

```
Out[41]: Investigation.Type      category
Location                       object
Abbreviation                   object
Country                       object
Aircraft.damage                category
Make                          object
Model                         object
Number.of.Engines              object
Engine.Type                   category
Purpose.of.flight              category
Total.Fatal.Injuries           float64
Total.Serious.Injuries         float64
Total.Minor.Injuries           float64
Total.Uninjured                float64
Weather.Condition              category
dtype: object
```

Exporting the cleaned dataset to a new .csv file

```
In [42]: # Save the cleaned DataFrame to a .csv file
# Set index = False to prevent pandas from creating a redundant index column
df.to_csv("data/cleaned-aviation-data.csv", index=False, encoding='latin1')
```

## Data Modeling

Loading the .csv file of the cleaned data

```
In [43]: # Load the cleaned .csv file and creating a new dataframe
df_clean = pd.read_csv("data/cleaned-aviation-data.csv", encoding='latin1', low_memory=False)
df_clean.head()
```

Out[43]:

	Investigation.Type	Location	Abbreviation	Country	Aircraft.damage	Make	Model	Number.ofEngines
0	Accident	FILLMORE	UT	United States	Substantial	BEECH	K35	1.0
1	Accident	BRAWLEY	CA	United States	Substantial	BELL	OH-58C	1.0
2	Accident	SOMERSET	KY	United States	Destroyed	BEECH	C-90	2.0
3	Accident	PLAINVILLE	CT	United States	Substantial	CESSNA	T310R	2.0
4	Accident	RAYVILLE	LA	United States	Substantial	AIR TRACTOR	AT-401	1.0

```
In [44]: # Examine the columns of the df_clean DataFrame
df_clean.columns
```

Out[44]: Index(['Investigation.Type', 'Location', 'Abbreviation', 'Country', 'Aircraft.damage', 'Make', 'Model', 'Number.ofEngines', 'Engine.Type', 'Purpose.of.flight', 'Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition'], dtype='object')

```
In [45]: # Print the shape of df_clean
df_clean.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 2534 rows  
This data set consists of 15 columns

## The Least Safe Aircraft

To gain insight on the least safe aircrafts, I group the `Model` and the `Make` variable and plot a barplot against `Total.Fatal.Injuries`

```
In [46]: # Group by 'Make' and 'Model', and sum 'Total.Fatal.Injuries'
Fatality_by_make_model = df_clean.groupby(['Make', 'Model'])['Total.Fatal.Injuries']

# Create a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in Fatality_by_make_model]

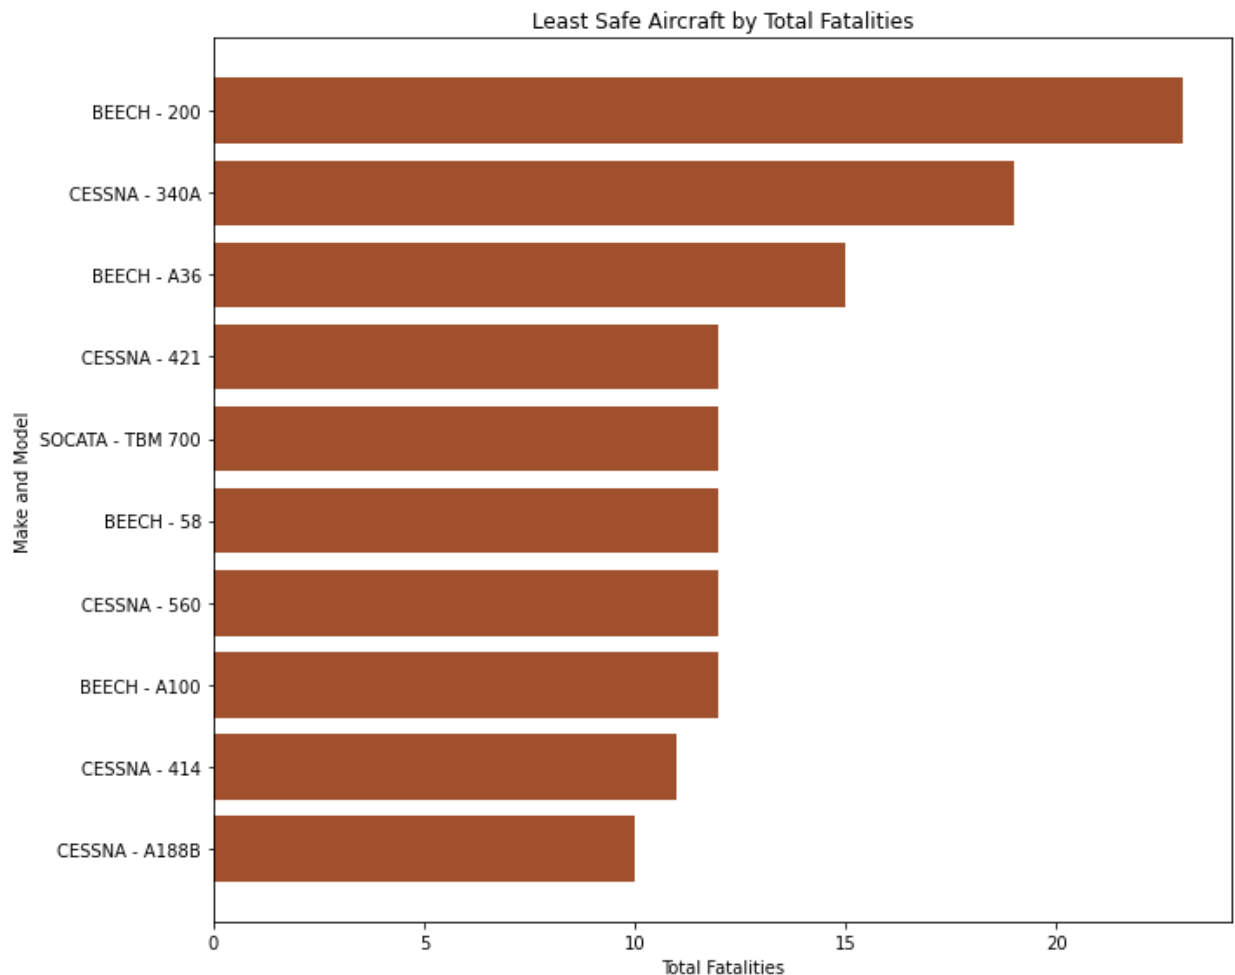
# Create a horizontal bar plot using Matplotlib
fig, ax = plt.subplots(figsize=(10, 8))

# Reverse the order in which bars are plotted to descending
ax.barh(make_model_labels[::-1], Fatality_by_make_model.values[::-1], color='sien')

# Set and customize the plot's Title, X and Y labels
ax.set_title('Least Safe Aircraft by Total Fatalities')
ax.set_xlabel('Total Fatalities')
ax.set_ylabel('Make and Model')

fig.tight_layout()

# Save the plot to the image folder
plt.savefig("./images/least-safe-aircraft.png", dpi=300, facecolor='white')
plt.show()
```





## The Most Safe Aircraft

To gain insight on the safest aircraft model and make, I group the `Model` and the `Make` variable and plot a barplot against `Total.Uninjured`

```
In [47]: # Group by 'Make' and 'Model', and sum 'Total.Uninjured'
uninjured_by_make_model = df_clean.groupby(['Make', 'Model'])['Total.Uninjured'].sum

# Create a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in uninjured_by_make_model]


# Plot a horizontal bar plot using Matplotlib
fig, ax = plt.subplots(figsize=(10, 8))

# Reverse the order in which bars are plotted to descending
ax.barh(make_model_labels[::-1], uninjured_by_make_model.values[::-1], color='royalblue')

# Set and customize the plot's Title, X and Y labels
ax.set_title('Most Safe Aircraft Total Uninjured')
ax.set_xlabel('Total Uninjured Passengers')
ax.set_ylabel('Make and Model')

fig.tight_layout()

# Save the plot to the image folder
plt.savefig("./images/most-safe-aircraft.png", dpi=300, facecolor='white')
plt.show()
```



To determine the safest aircraft models for each of the three civil aviation services the company can venture into; the three categorical values for the `Purpose.of.Flight` columns are plotted in a barplot against uninjured passengers `Total.Uninjured`.

## Recommended Aircraft for Targeted Aviation Services

```
In [48]: # Filter for the relevant Purpose.of.flight values
df_filtered = df_clean[df_clean['Purpose.of.flight'].isin(['Aerial Application', 'B

# Group by 'Purpose.of.flight', 'Make', and 'Model', and sum 'Total.Uninjured'
uninjured_by_purpose_make_model = df_filtered.groupby(['Purpose.of.flight', 'Make',

# Find the safest aircraft (highest 'Total.Uninjured') for each purpose
safest_aircraft = uninjured_by_purpose_make_model.loc[uninjured_by_purpose_make_mod

# Create the "Make - Model" column
safest_aircraft['Make - Model'] = safest_aircraft['Make'] + ' - ' + safest_aircraft

# Create a bar plot using Matplotlib with subplots
fig, ax = plt.subplots(figsize=(12, 8))

# Define colors assigned to bars
colors = ['tab:blue', 'tab:orange', 'tab:green']


# Plot the horizontal barplot with individual labels for each bar
bars = ax.bar(safest_aircraft['Purpose.of.flight'], safest_aircraft['Total.Uninjure

# Add a legend
legend_labels = safest_aircraft['Make - Model'].tolist()
ax.legend(bars, legend_labels, title="Safest Aircraft")

# Set and customize the plot's Title, X and Y labels
ax.set_title('Recommended Aircraft')
ax.set_xlabel('Purpose of Flight')
ax.set_ylabel('Total Uninjured Passengers')

fig.tight_layout()

# Save the plot to the image folder
plt.savefig("./images/recommended-aircraft.png", dpi=300, facecolor='white')
plt.show()
```



## Evaluation

The baseline model shed the following insights:

- The BEECH-200 (8-seater), the CESSNA-340A (6-seater), and the BEECH-A36 (6-seater), are the top-three most risky aircraft overall.
- The CESSNA-A188B (1-seater), the BELL-206B (5-seater), and the GRUMMAN-G-164A (1-seater) are the top-three safest aircraft models overall.

### Recommendations:

- The CESSNA-560XL (10-seater) is the safest aircraft for Executive/corporate flights.
- The DEHAVILLAND-DHC-8-311 (50-seater) is the safest airplane for business flights.
- The CESSNA-A188B (1-seater) is the safest aircraft for Aerial Applications.

The number of engines for these aircraft models.

- BEECH-200: 2 engines
- CESSNA-340A: 2 engines
- BEECH-A36: 1 engine
- BELL-206B: 1 engine
- GRUMMAN-G-164A: 1 engine
- CESSNA-560XL: 2 engines
- CESSNA-A188B: 1 engine
- DEHAVILLAND-DHC-8-311: 2 engines

Multi-engine aircraft are typically safer in comparison to single-engine airplanes (Pilot Institute, 2023). More than one engine avails redundancy to propulsion units. Pilots can recalibrate controls and continue flying if one engine stalls or malfunctions. In contrast, such an incident in a single-engine aircraft poses a substantial risk of

crashing if the aliding distance falls short of a runway or a relatively flat surface for an emergency landing. Thus.

```
In [49]: # Create a copy of df_clean to manipulate
df_modified = df_clean.copy()

# Check the number of rows and columns of the DataFrame
df_modified.shape
print(f"This data set consists of {df.shape[0]} rows")
print(f"This data set consists of {df.shape[1]} columns")
```

This data set consists of 2534 rows  
This data set consists of 15 columns

```
In [50]: # Apply a lambda function to drop entries for single-engine aircraft
df_modified = df_modified[df_modified['Number.of.Engines'].apply(lambda x: x >= 2)]
```

```
In [51]: # Confirm if single-engine entries dropped
df_modified['Number.of.Engines'].value_counts()
```

```
Out[51]: 2.0    369
3.0      7
4.0      3
Name: Number.of.Engines, dtype: int64
```

Replotting bar plots to determine the most risky, the safest, and the appropriate airplanes (that have more than one engine) for the company's operations

## The Least Safe Multi-Engine Aircraft

```
In [52]: # Group by 'Make' and 'Model', and sum 'Total.Fatal.Injuries'
Fatality_by_make_model = df_modified.groupby(['Make', 'Model'])['Total.Fatal.Injuries'].sum()

# Create a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in Fatality_by_make_model.index]

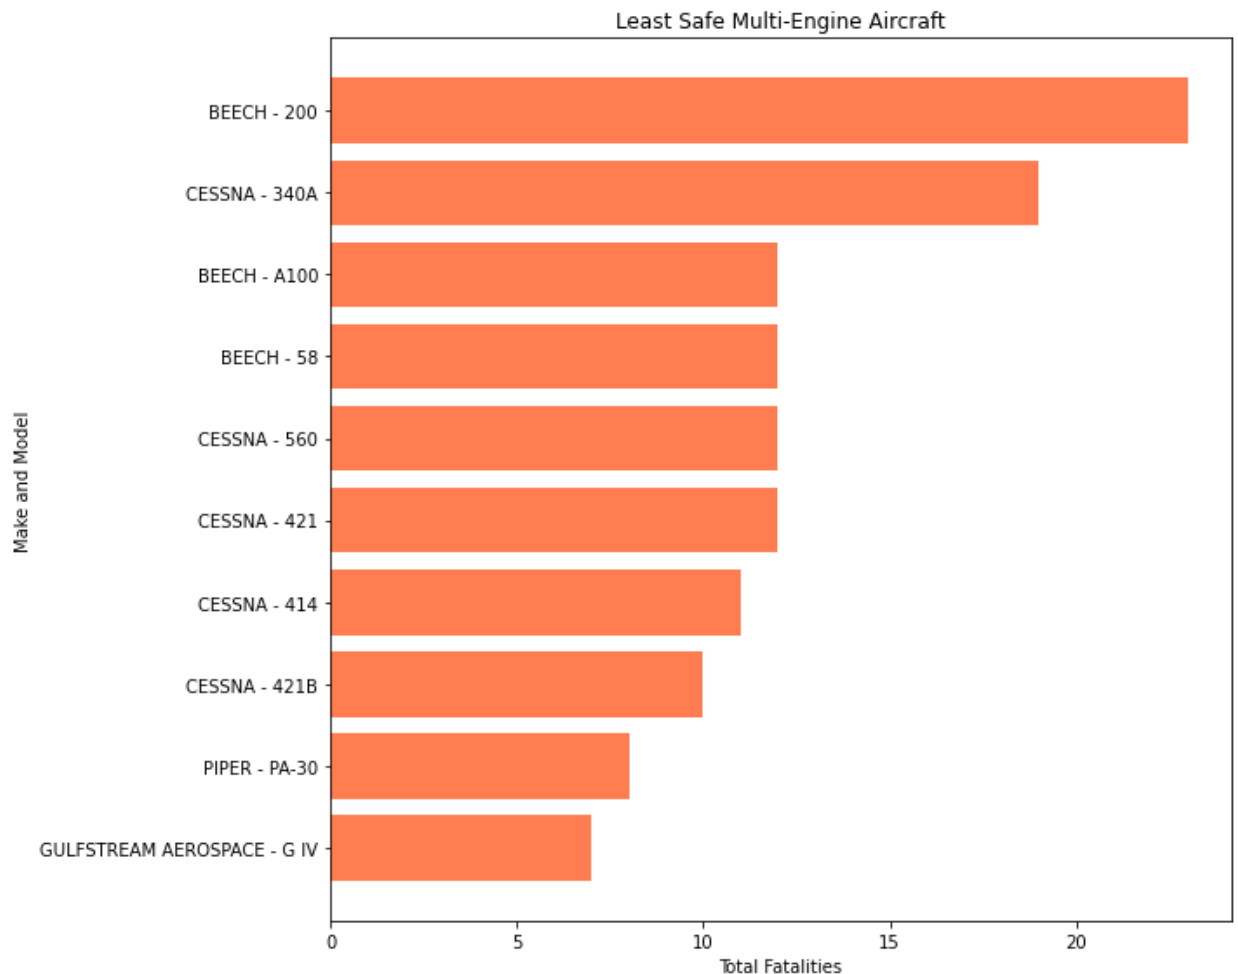
# Create a horizontal bar plot using Matplotlib
fig, ax = plt.subplots(figsize=(10, 8))

# Reverse the order in which bars are plotted to descending
ax.barh(make_model_labels[::-1], Fatality_by_make_model.values[::-1], color='coral')

# Set and customize the plot's Title, X and Y labels
ax.set_title('Least Safe Multi-Engine Aircraft')
ax.set_xlabel('Total Fatalities')
ax.set_ylabel('Make and Model')

fig.tight_layout()

# Save the plot to the image folder
plt.savefig("./images/least-safe-multi-engine-aircraft.png", dpi=300, facecolor='white')
plt.show()
```



## The Most Safe Multi-Engine Aircraft

```
In [53]: # Group by 'Make' and 'Model', and sum 'Total.Uninjured'
uninjured_by_make_model = df_modified.groupby(['Make', 'Model'])['Total.Uninjured']

# Create a list of the Make and Model labels for the y-axis
make_model_labels = [f"{make} - {model}" for make, model in uninjured_by_make_model]

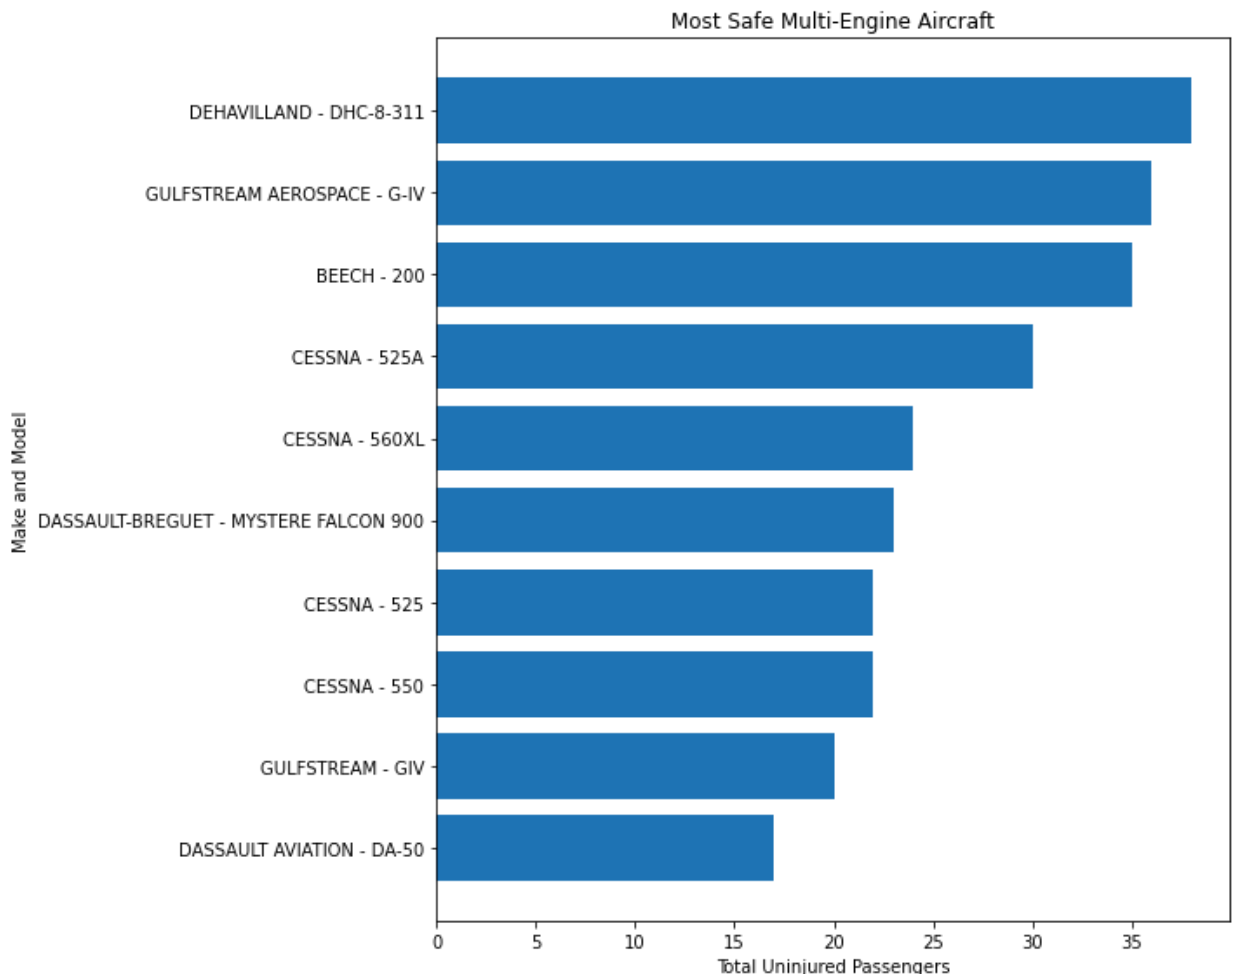
# Create a horizontal bar plot using Matplotlib
fig, ax = plt.subplots(figsize=(10, 8))

# Reverse the order in which bars are plotted to descending
ax.barh(make_model_labels[::-1], uninjured_by_make_model.values[::-1])

# Set and customize the plot's Title, X and Y labels
ax.set_title('Most Safe Multi-Engine Aircraft')
ax.set_xlabel('Total Uninjured Passengers')
ax.set_ylabel('Make and Model')

fig.tight_layout()

# Save to the image folder
plt.savefig("./images/most-safe-multi-engine-aircraft.png", dpi=300, facecolor='palegoldenrod')
plt.show()
```



## Recommended Multi-Engine Aircraft

```
In [54]: # Filter for the relevant Purpose.of.flight values
df_filtered = df_modified[df_modified ['Purpose.of.flight'].isin(['Aerial Application', 'Cargo', 'Passenger'])]

# Group by 'Purpose.of.flight', 'Make', and 'Model', and sum 'Total.Uninjured'
uninjured_by_purpose_make_model = df_filtered.groupby(['Purpose.of.flight', 'Make', 'Model']).sum()

# Find the safest aircraft (highest 'Total.Uninjured') for each purpose
safest_aircraft = uninjured_by_purpose_make_model.loc[uninjured_by_purpose_make_model.groupby('Purpose.of.flight')['Total.Uninjured'].idxmax()]

# Create the "Make - Model" column
safest_aircraft['Make - Model'] = safest_aircraft['Make'] + ' - ' + safest_aircraft['Model']

# Plot a bar plot using Matplotlib with subplots
fig, ax = plt.subplots(figsize=(12, 8))

# Define colors assigned to bars
colors = ['tab:blue', 'tab:orange', 'tab:green']

# Plot the horizontal barplot with individual labels for each bar
bars = ax.bar(safest_aircraft['Purpose.of.flight'], safest_aircraft['Total.Uninjured'], color=colors)

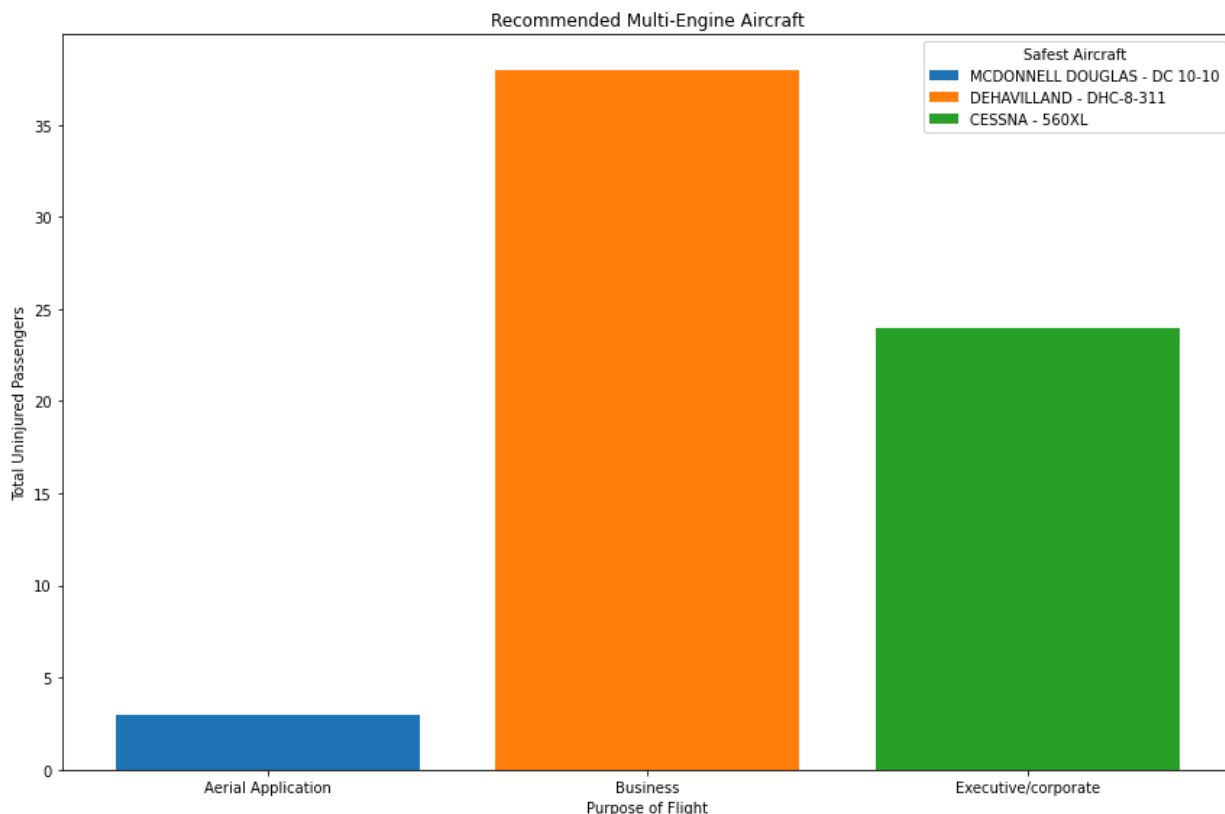
# Add a legend
legend_labels = safest_aircraft['Make - Model'].tolist()
ax.legend(bars, legend_labels, title="Safest Aircraft")

# Set and customize the plot's Title, X and Y labels
ax.set_title('Recommended Multi-Engine Aircraft')
ax.set_xlabel('Purpose of Flight')
ax.set_ylabel('Total Uninjured Passengers')

fig.tight_layout()

# Save the plot to the image folder
plt.savefig("../images/recommended-multi-engine-aircraft.png", dpi=300, facecolor='white')
plt.show()
```





The modified model sheds the following insight:

- The BEECH-200 (13-seater), the CESSNA-340A (6-seater), and the BEECH-A400 (8-seater) are the top-three most risky multi-engine aircraft to operate.
- The DEHAVILLAND DHC-8-311 (50-seater), the GULFSTREAM AEROSPACE-G-IV (10-seater), and the BEECH-200 (13-seater) are the top-three safest multi-engine aircraft to operate.

#### Recommendations:

- The CESSNA 560XL(10-seater) is the safest multi-engine aircraft for Executive/corporate flights.
- The DEHAVILLAND -DHC-8-311(50-seater) is the safest multi-engine aircraft for business flights.
- The MCDONNELL DOUGLAS-DC 10-10(250-seater) is the safest multi-engine aircraft for Aerial Applications.

The number of engines for these aircraft models.

- BEECH-200: 2 engines
- CESSNA-340A: 2 engines
- BEECH-A400: 2 engines
- GULFSTREAM AEROSPACE-G-IV: 2 engines
- CESSNA 560XL: 2 engines
- DEHAVILLAND DHC-8-311: 2 engines
- MCDONNELL DOUGLAS -DC 10-10: 3 engines

The findings from the results yielded by the baseline model are compared to those yielded by the modified model. After entering the industry, the recommended safest aircraft model for the aviation services the company will offer is based on respective applicability and operational logistics.

## Conclusions

The analysis yields three recommendations on the aircraft models the company should procure and operate after entering the commercial aviation industry.

- **The CESSNA-560XL aircraft is recommended for Executive/ Corporate flights:** The baseline and modified model conform the aircraft is safest for executive and corporate flights.
- **The DEHAVILLAND DHC-8-311 aircraft is recommended for Business Flights:** The baseline and modified model conform the aircraft is safest for business flights.
- **The CESSNA-A188B aircraft is recommended for Aerial Applications:** The modified model proposes the MCDONNELL DOUGLAS-DC 10-10 (a three-engine, 250-seater) aircraft for aerial applications. The proposed alternative by the modified model is rejected because aerial applications typically include agricultural activities such as spraying crop fields. Hence, the single-engine CESSNA-A188B is recommended for aerial applications.

## References

Altundag, A., & Wynn, M. (2024). Advanced Analytics and Data Management in the Procurement Function: An Aviation Industry Case Study. *Electronics*, 13(8), 1554. <https://doi.org/10.3390/electronics13081554> (<https://doi.org/10.3390/electronics13081554>).

Boyd, D. D., & Stolzer, A. (2016). Accident-precipitating factors for crashes in turbine-powered general aviation aircraft. *Accident Analysis & Prevention*, 86, 209-216. <https://doi.org/10.1016/j.aap.2015.10.024> (<https://doi.org/10.1016/j.aap.2015.10.024>).

NTSB. (2023). Aviation Accident Database & Synopses, up to 2023. <https://doi.org/10.34740/KAGGLE/DSV/4875576> (<https://doi.org/10.34740/KAGGLE/DSV/4875576>).

Luo, J. (2022). Data-driven innovation: What is it?. *IEEE Transactions on Engineering Management*, 70(2), 784-790. <https://doi.org/10.1109/TEM.2022.3145231> (<https://doi.org/10.1109/TEM.2022.3145231>).

Pilot Institute. (2023). Single Engine Vs. Multi Engine: Which is Better? <https://pilotinstitute.com/single-vs-multi-engine/> (<https://pilotinstitute.com/single-vs-multi-engine/>).